

Image Classification And Object Detection

For researchers interested in studying Earth science with deep learning.

All resources in lectures are available at https://github.com/MrXiaoXiao/DLiES

Deep Learning in Earth Science Lecture 2 By Xiao Zhuowei





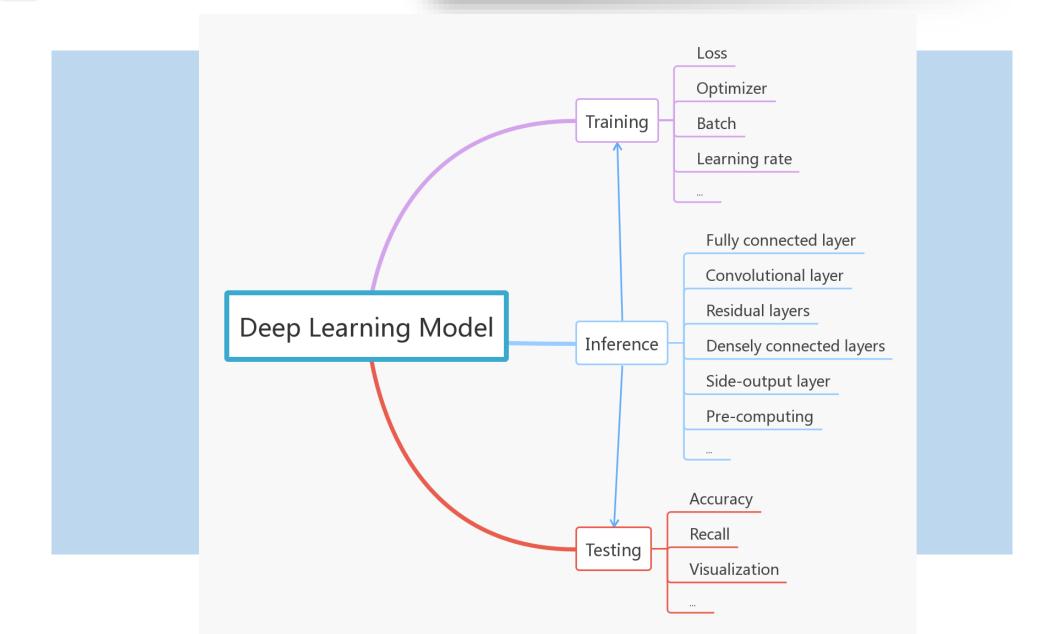
Classification with Convolutional Neural Networks



Object Detection with Bounding Box



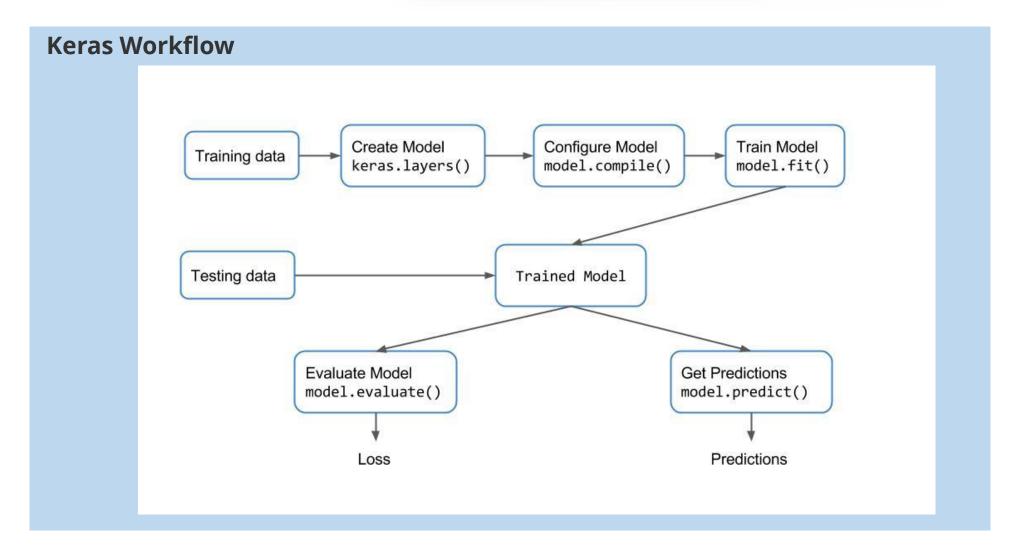
Discussions



Keras

Keras is a high-level neural networks API, written in Python and capable of running on top of <u>TensorFlow</u>, <u>CNTK</u>, or <u>Theano</u>. It was developed with a focus on enabling fast experimentation.

Being able to go from idea to result with the least possible delay is key to doing good research.



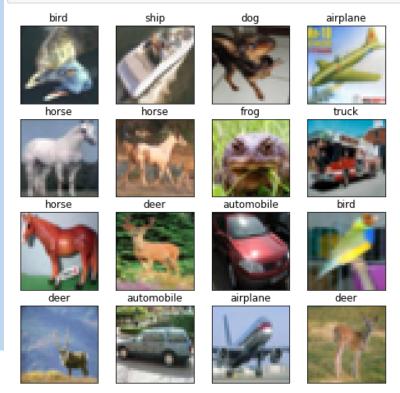


Getting started with the Keras

Prepare dataset

```
import tensorflow as tf
from tensorflow.keras.datasets import cifar10
#train and test using cifar10 dataset
(x_train, y_train), (x_test, y_test) = cifar10.load_data()
print('x train shape:', x train.shape)
print(x_train.shape[0], 'train samples')
print(x test.shape[0], 'test samples')
#Preprocessing
x_train = x_train.astype('float32')
x test = x test.astype('float32')
x train /= 255
x test /= 255
num classes = 10
# Convert class vectors to binary class matrices.
y train = tf.keras.utils.to categorical(y train, num classes)
y test = tf.keras.utils.to categorical(y test, num classes)
/home/wangi/anaconda3/envs/DLiES/lib/python3.6/importlib/ bootstrap.py:219: RuntimeWarning: numpy.dtype size changed, ma
y indicate binary incompatibility. Expected 88 from C header, got 96 from PyObject
 return f(*args, **kwds)
x_train shape: (50000, 32, 32, 3)
50000 train samples
10000 test samples
```

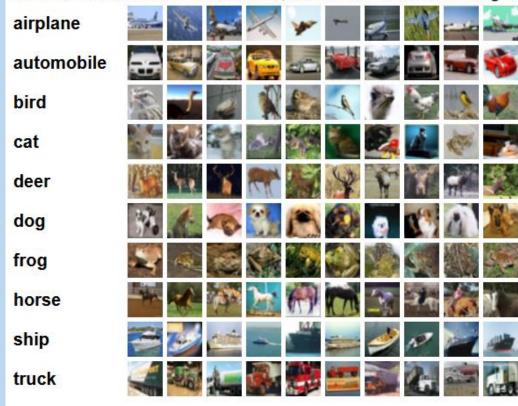
Check dataset



The CIFAR-10 dataset

The CIFAR-10 dataset consists of 60000 32x32 colour images in 10 classes, with 6000 images per class. There are 50000 training images and 10000 test images.

http://www.cs.toronto.edu/ ~kriz/cifar.html Here are the classes in the dataset, as well as 10 random images from each:



Getting started with the Keras

Create a model In [4]: #The Sequential model is a linear stack of layers. model = tf.keras.Sequential() Add layers In [5]: from tensorflow.keras.layers import Dense, Dropout, Activation, Flatten from tensorflow.keras.layers import Conv2D, MaxPooling2D #You can also simply add layers via the .add() method model.add(Conv2D(32, (3, 3), padding='same',input_shape=x_train.shape[1:], activation='relu')) model.add(Conv2D(32, (3, 3), activation='relu')) model.add(MaxPooling2D(pool_size=(2, 2))) model.add(Dropout(0.25)) model.add(Conv2D(64, (3, 3), padding='same', activation='relu')) model.add(Activation('relu')) model.add(Conv2D(64, (3, 3), padding='same', activation='relu')) model.add(MaxPooling2D(pool_size=(2, 2))) model.add(Dropout(0.25)) model.add(Flatten()) model. add (Dense (512, activation='relu')) model.add(Dropout(0.5)) model. add (Dense (10, activation='softmax'))

Getting started with the Keras

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 32)	896
conv2d_1 (Conv2D)	(None, 32, 32, 32)	9248
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
dropout (Dropout)	(None, 16, 16, 32)	0
conv2d_2 (Conv2D)	(None, 16, 16, 64)	18496
activation (Activation)	(None, 16, 16, 64)	0
conv2d_3 (Conv2D)	(None, 16, 16, 64)	36928
max_pooling2d_1 (MaxPooling2	(None, 8, 8, 64)	0

Initiate optimizer

```
In [7]: opt = tf.keras.optimizers.Adam()
```

Configure model

Train model

```
In [9]: #Keras models are trained on Numpy arrays of input data and labels.
     #For training a model, you will typically use the fit function.
     history = model.fit(x_train, y_train,
                   batch size=32,
                   epochs=10,
                   shuffle=True)
      Epoch 1/10
      Epoch 2/10
      Epoch 3/10
     50000/50000 [============= ] - 14s 283us/step - loss: 0.9782 - acc: 0.6532
      Epoch 4/10
     50000/50000 [============= ] - 14s 277us/step - loss: 0.8747 - acc: 0.6925
      Epoch 5/10
      Epoch 6/10
     50000/50000 [===========] - 14s 277us/step - loss: 0.7452 - acc: 0.7386
      Epoch 7/10
     50000/50000 [=============] - 14s 278us/step - loss: 0.7060 - acc: 0.7504
      Epoch 8/10
     50000/50000 [=======] - 14s 276us/step - loss: 0.6683 - acc: 0.7654
      Epoch 9/10
      50000/50000 [============= ] - 13s 266us/step - loss: 0.6360 - acc: 0.7778
      Epoch 10/10
      50000/50000 [============= ] - 14s 273us/step - loss: 0.6098 - acc: 0.7857
```

Test model

Test accuracy: 0.7711

Check predictions

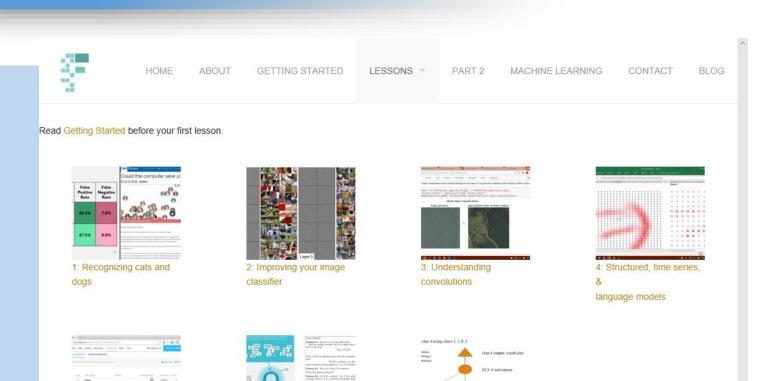
```
In [11]: preds = model.predict(x_test)
         plt.figure(figsize=(8,8))
          for i in range(16):
             plt.subplot(4, 4, 1 + i, xticks=[], yticks=[])
             img_id = np.random.randint(10000)
             im = x_test[img_id,::]
             plt.title(class_names[preds[img_id].argmax()])
             plt.imshow(im)
         plt.show()
                           automobile
                                             bird
                                                           horse
                           automobile
                                                         automobile
```



https://www.fast.ai/

Inside the training loop

If you can code, you can do deep learning.



RNNs from scratch

7: Resnets from scratch





Classification with Convolutional Neural Networks



Object Detection with Bounding Box



Discussions

DeepSat (SAT-6) Airborne Dataset
405,000 image patches each of size
28x28 and covering 6 landcover
classes

https://www.kaggle.com/crawford/deepsat-sat6





Kaggle is an online community of data scientists and machine learners, owned by Google, Inc. Kaggle allows users to find and publish data sets, explore and build models in a web-based data-science environment, work with other data scientists and machine learning engineers, and enter competitions to solve data science challenges

Classifying Satellite Images

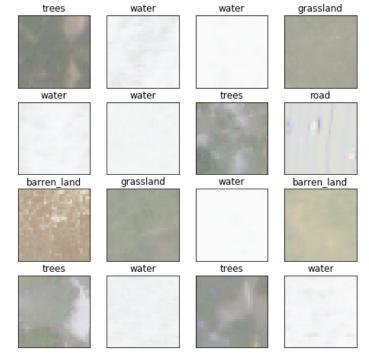
Prepare dataset

```
#require panadasy
import pandas as pd # data processing, CSV file I/O (e.g. pd. read csv)
import numpy as np
# Method to load data and images
def load data and labels (data, labels):
    data df = pd. read csv(data, header=None)
   X = data_df.values.reshape((-1, 28, 28, 4)).clip(0, 255).astype(np.uint8)
    labels_df = pd. read_csv(labels, header=None)
   Y = labels_df.values.getfield(dtype=np.int8)
    return X, Y
data dir = 'F:/deepsat sat6'
x_train, y_train = load_data_and_labels(data=' {} /X_train_sat6.csv'.format(data_dir),
                                       labels='{}/v train sat6.csv'.format(data dir))
x_test, y_test = load_data_and_labels(data='{}/X_test_sat6.csv'.format(data_dir),
                                     labels='{}/v test sat6.csv'.format(data dir))
print(pd.read_csv('{})/sat6annotations.csv'.format(data_dir), header=None))
# Print shape of all training, testing data and labels
# Labels are already loaded in one-hot encoded format
print('x_train_shape : {}'.format(x_train.shape)) # (324000, 28, 28, 4)
print('y_train_shape : {}'.format(y_train.shape)) # (324000, 6)
print('x_test_shape: {}'.format(x_test.shape)) # (81000, 28, 28, 4)
print('y_test_shape : {}'.format(y_test.shape)) # (81000, 6)
             0 1 2 3 4 5 6
     building 1 0 0 0 0
1 barren_land 0 1 0 0 0 0
     grassland 0 0 0 1 0 0
         road 0 0 0 0 1 0
        water 0 0 0 0 0 1
x_train_shape : (324000, 28, 28, 4)
y_train_shape : (324000, 6)
x_test_shape: (81000, 28, 28, 4)
v test_shape : (81000, 6)
```

Classification with Convolutional Neural Networks

Classifying Satellite Images

Check dataset



Classifying Satellite Images

Create model

```
import tensorflow as tf
from tensorflow.keras.layers import Conv2D, MaxPool2D, Dense, Dropout, Flatten
from tensorflow.keras.models import Sequential
#Create model
model = Sequential()
#Add layers
model.add(Conv2D(16, (3,3), activation='relu', input shape=(28,28,4)))
model.add(Conv2D(32, (3,3), activation='relu'))
model.add(MaxPool2D(pool size=(2,2)))
model.add(Dropout(0.5))
model.add(Conv2D(32, (3,3), activation='relu'))
model.add(Conv2D(64, (3,3), activation='relu'))
model.add(MaxPool2D(pool size=(2,2)))
model.add(Dropout(0.5))
model.add(Flatten())
model.add(Dense(12B, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(6, activation='softmax'))
#configure model
model.compile(loss='categorical crossentropy', optimizer='adam', metrics=['accuracy'])
print(model.summary())
```

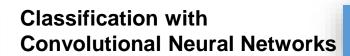


Classifying Satellite Images

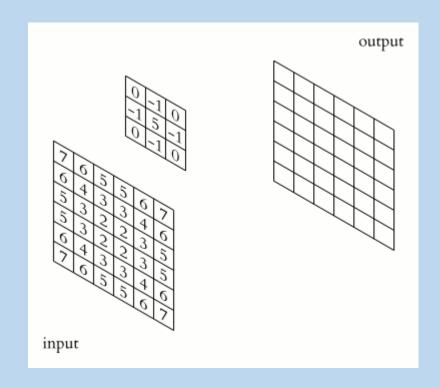
Layer (type)	Output	Shape	Param #
conv2d (Conv2D)	(None,	26, 26, 16)	592
conv2d_1 (Conv2D)	(None,	24, 24, 32)	4640
max_pooling2d (MaxPooling2D)	(None,	12, 12, 32)	0
dropout (Dropout)	(None,	12, 12, 32)	0
conv2d_2 (Conv2D)	(None,	10, 10, 32)	9248
conv2d_3 (Conv2D)	(None,	8, 8, 64)	18496
max_pooling2d_1 (MaxPooling2	(None,	4, 4, 64)	0
dropout_1 (Dropout)	(None,	4, 4, 64)	0
flatten (Flatten)	(None,	1024)	0
dense (Dense)	(None,	128)	131200
dropout_2 (Dropout)	(None,	128)	0
dense_1 (Dense)	(None,	6)	774
T-4-1 164 050			

Total params: 164,950 Trainable params: 164,950 Non-trainable params: 0

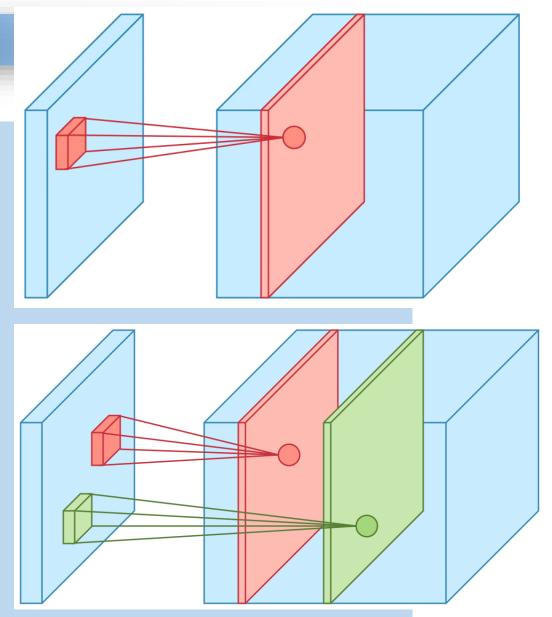
None



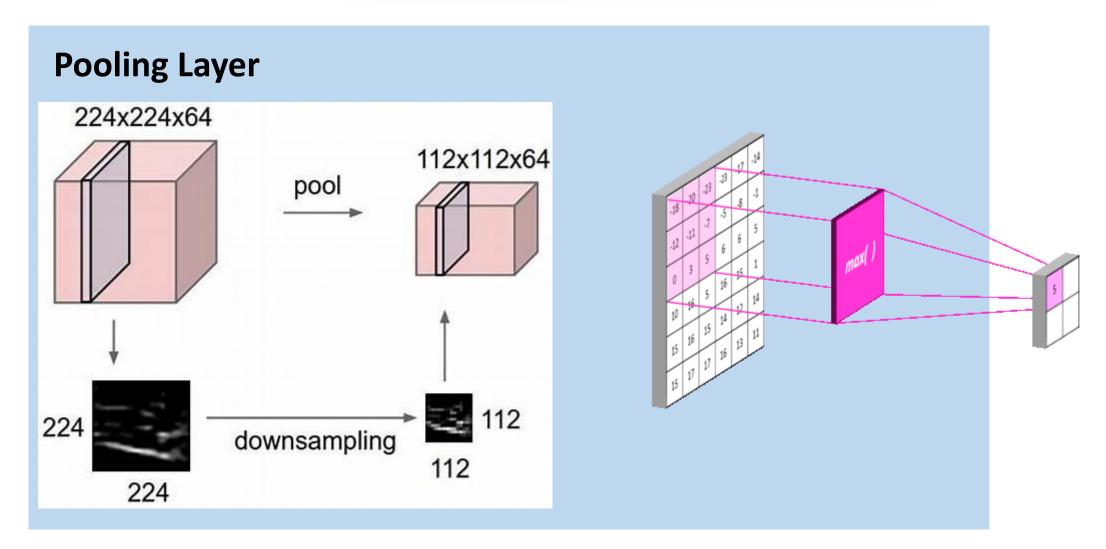
Convolutional Layer







https://towardsdatascience.com/applied-deep-learning-part-4-convolutional-neural-networks-584bc134c1e2



Dropout

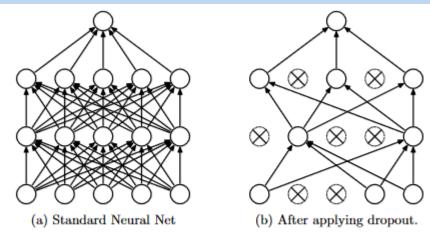
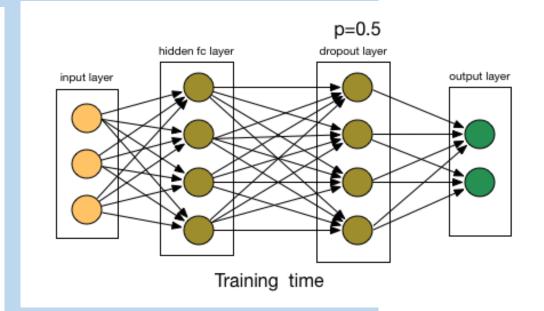


Figure 1: Dropout Neural Net Model. Left: A standard neural net with 2 hidden layers. Right: An example of a thinned net produced by applying dropout to the network on the left. Crossed units have been dropped.

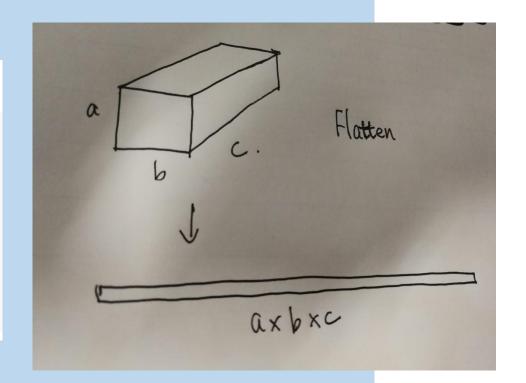


(Srivastava et al., 2014)

https://chatbotslife.com/regularization-in-deep-learning-f649a45d6e0

Flatten Layer

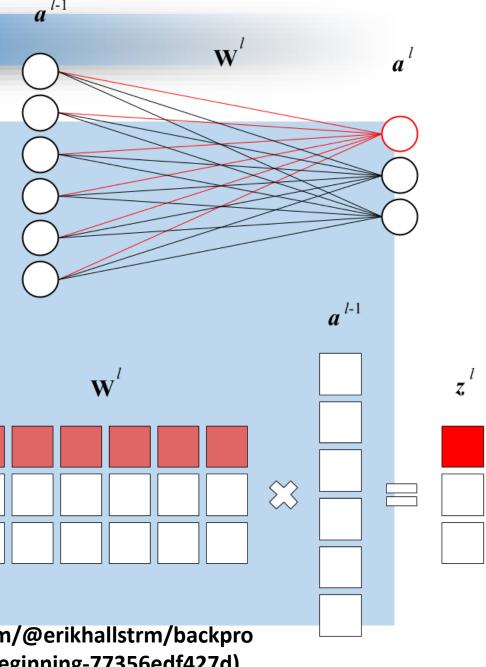
Example:



https://www.tensorflow.org/api_docs/pytho n/tf/keras/layers/Flatten

Densely(Fully) Connected Layer

$$\mathbf{a}^l = \sigma(\mathbf{W}^l \mathbf{a}^{l-1} + \mathbf{b}^l)$$



Images From

(https://medium.com/@erikhallstrm/backpro pagation-from-the-beginning-77356edf427d)

Classifying Satellite Images

Train model

Test model

Check predictions

Classifying Satellite Images

```
preds = model.predict(x test)
plt.figure(figsize=(8,8))
for i in range(16):
    plt.subplot(4, 4, 1 + i, xticks=[], yticks=[])
    img_id = np.random.randint(81000)
    im = x_test[img_id,::]
    plt.title(class_names[preds[img_id].argmax()])
    plt.imshow(im)
plt.show()
                                                  grassland
     water
                    trees
                                    water
     water
                                  grassland
                    water
                                                    water
                                 barren land
                                                    road
     water
                    water
   grassland
                    water
                                    water
                                                    water
```





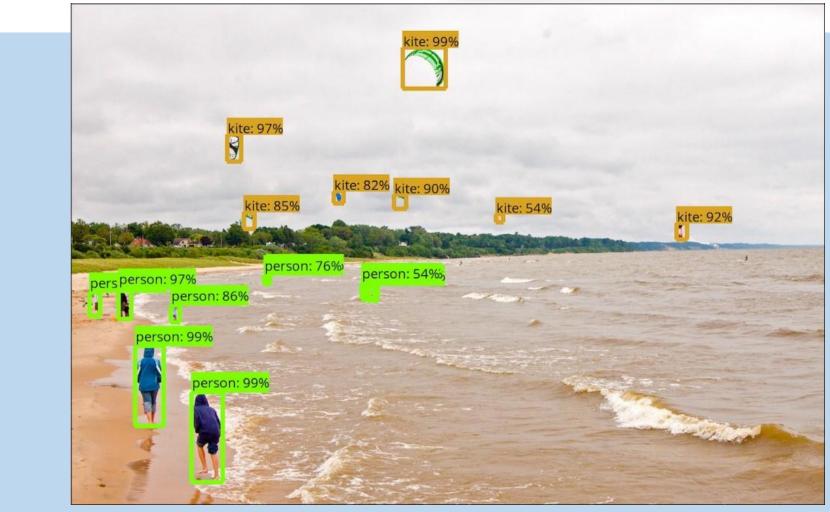
Classification with Convolutional Neural Networks



Object Detection with Bounding Box

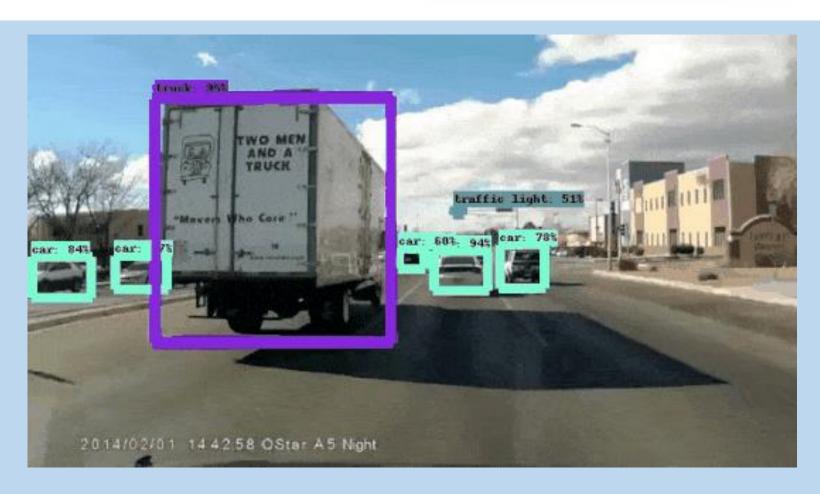


Discussions



https://towardsdatascience.com/is-google-tensorflow-object-

detection-api-the-easiest-way-to-implement-image-recognition-a8bd1f500ea0

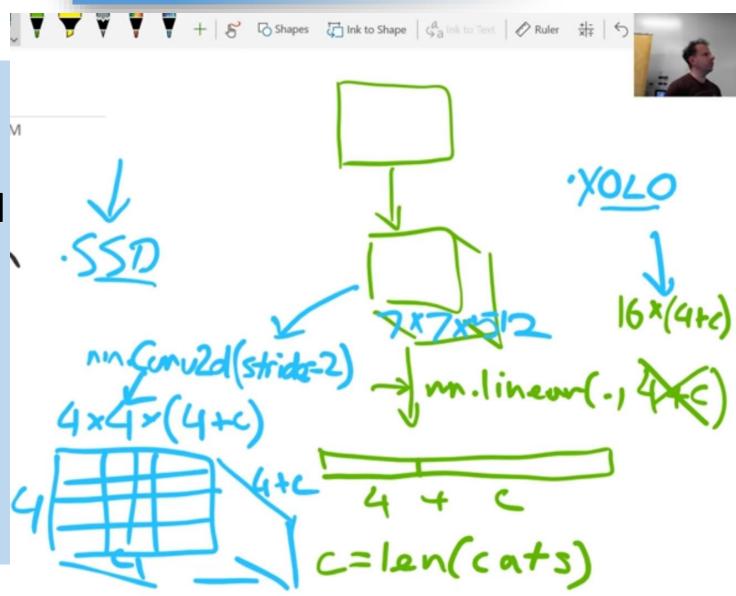


https://towardsdatascience.com/is-google-tensorflow-object-detection-api-the-easiest-way-to-implement-image-recognition-a8bd1f500ea0



fast.ai courses strongly recommended

https://course.fast.ai/lessons/lesson8.html https://course.fast.ai/lessons/lesson9.html



Bounding Box Example

In geometry, the minimum or smallest bounding or enclosing box for a point set (S) in N dimensions is the box with the smallest measure (area, volume, or hypervolume in higher dimensions) within which all the points lie.

(Wikipedia)



http://host.robots.ox.ac.uk/pascal/VOC/voc2006/examples/index.html

SPPNet: Spatial Pyramid Pooling in Deep Convolutional

Networks for Visual Recognition

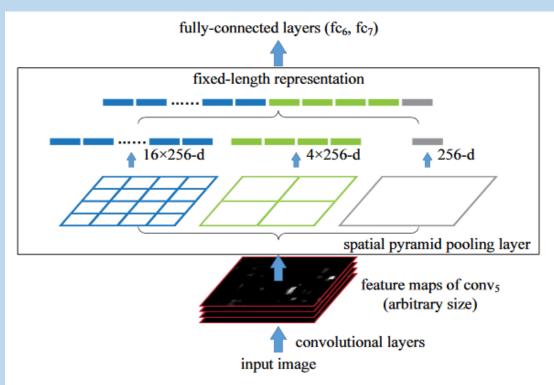
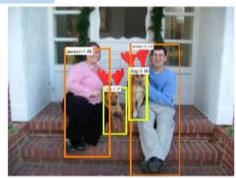
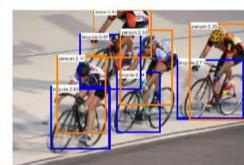


Figure 3: A network structure with a **spatial pyramid pooling layer**. Here 256 is the filter number of the $conv_5$ layer, and $conv_5$ is the last convolutional layer.



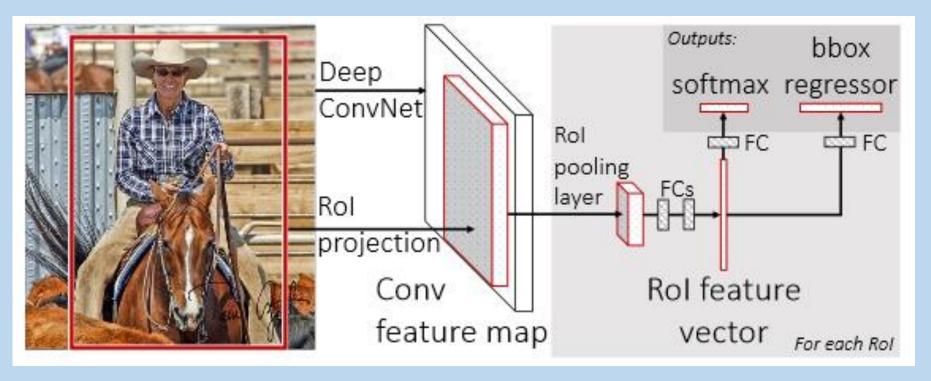






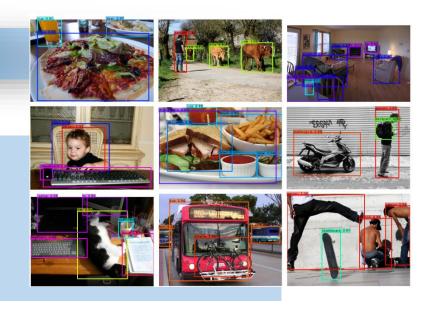
(He et al. 2014)

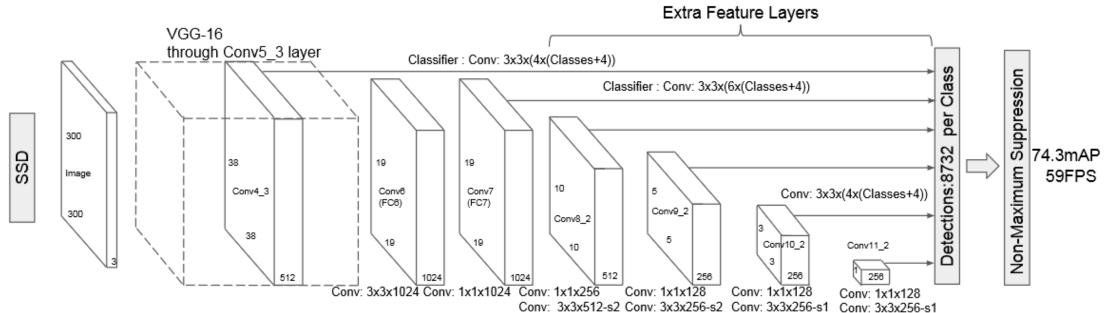
Fast-RCNN: Fast Region-based Convolutional Network method

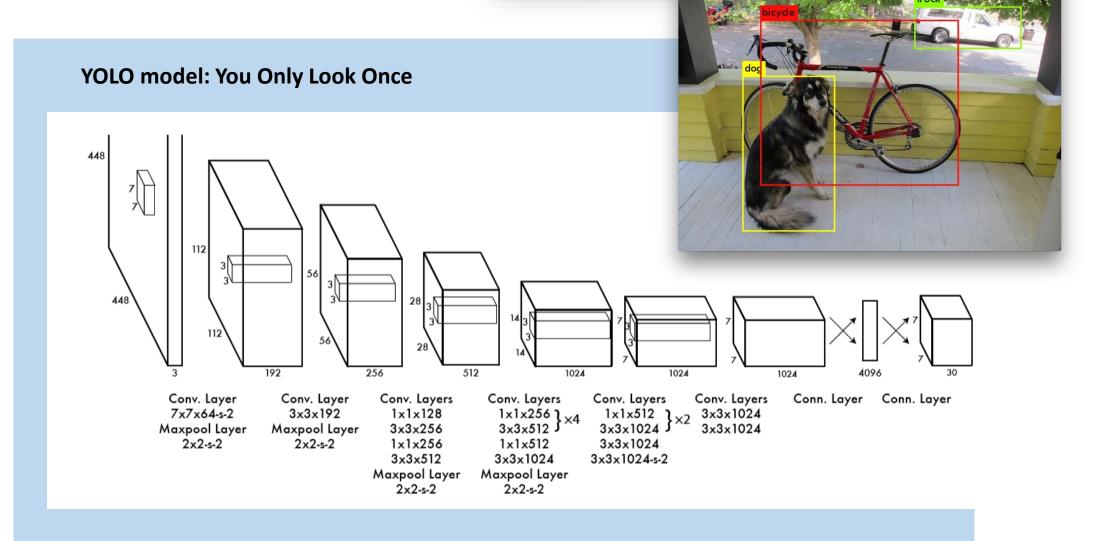


(Girshick, 2015)













Classification with Convolutional Neural Networks



Object Detection with Bounding Box



Discussions

Discussions



References

- [1] N. Srivastava, G. Hinton, A. Krizhevsky, I. Sutskever, and R. Salakhutdinov, "Dropout: A Simple Way to Prevent Neural Networks from Overfitting," *Journal of Machine Learning Research*, vol. 15, pp. 1929–1958, 2014.
- [2] K. He, X. Zhang, S. Ren, and J. Sun, "Spatial Pyramid Pooling in Deep Convolutional Networks for Visual Recognition," arXiv:1406.4729 [cs], vol. 8691, pp. 346–361, 2014.
- [3] R. Girshick, "Fast R-CNN," arXiv:1504.08083 [cs], Apr. 2015.
- [4] W. Liu *et al.*, "SSD: Single Shot MultiBox Detector," *arXiv:1512.02325* [cs], vol. 9905, pp. 21–37, 2016.
- [5] J. Redmon, S. Divvala, R. Girshick, and A. Farhadi, "You Only Look Once: Unified, Real-Time Object Detection," in *2016 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, Las Vegas, NV, USA, 2016, pp. 779–788.
- [6] J. Redmon and A. Farhadi, "YOLOv3: An Incremental Improvement," p. 6.